

High-Dimensional Data Analysis Project

Group 3H (Bagas Ari Wibawanto, Izzul Fattah Aji Pratama, Yusuf Romadhon)

2026-05-06

Table of contents

1	Introduction	1
2	Data and Preparation	1
3	Preliminary Analysis (EDA)	2
4	PCA Preparation	4
5	PCA and Latent Factors Analysis	5
5.1	Q1 Does the largest principal component coincide with market returns?	5
5.2	Q2 Do Market Return, SMB, and HML explain common variation in stock returns?	6
5.3	Q3 Which industries move in the same direction as the common factors?	7
5.4	Q4 Which industry has the highest systematic risk?	8
5.5	Q5 Which industry is counter-cyclical for hedging purposes?	9
6	Conclusion	10
7	Reference	11
7.1	AI Usage Statement	12
8	Appendix	12
8.1	Data and Preparation Code	12
8.2	Preliminary Analysis Code	14
8.3	PCA Preparation Code	17
8.4	Q1 Code	17
8.5	Q2 Code	17
8.6	Q3 Code	19
8.7	Q4 Code	20
8.8	Q5 Code	21

1 Introduction

The analysis of stock returns is important in finance because it helps explain the relationship between risk and return. This report analyses monthly returns of New York Stock Exchange (NYSE) from 83 companies from 2005 to 2019 (over 180 months) using high-dimensional data techniques. The study separates stock returns into systematic risk driven by overall market movements and idiosyncratic risk specific to individual firms or industries.

The objective is to investigate whether latent factors correspond to observable market benchmarks. Specifically, the report examines whether the first principal component represents the market, whether the Fama-French three-factor model explains common variation, which industries are most exposed to common factors, and whether any industry provides defensive or hedging characteristics.

2 Data and Preparation

Data preparation is the foundation of high-dimensional analysis. We utilize `SampleH.csv` (stock returns) and `MarketFactors.csv` (benchmarks).

Table 1: Sample of Dataset

Stock ID	Industry	Return	Market Return	SMB	HML
F85464	Wholesale Trade	-0.028	-2.76	-1.66	2.06
H89060	Finance, Insurance and Real Estate	-0.069	-2.76	-1.66	2.06
H87476	Finance, Insurance and Real Estate	-0.040	-2.76	-1.66	2.06
E79879	Transportation and Public Utilities	-0.322	-2.76	-1.66	2.06
H10777	Finance, Insurance and Real Estate	-0.042	-2.76	-1.66	2.06
B89901	Mining	0.154	-2.76	-1.66	2.06

The final dataset contains 14,940 observations across 9 variables with no missing values. Industry classifications include Mining (B), Construction (C), Manufacturing (D), Transportation and Public Utilities (E), Wholesale Trade (F), Retail Trade (G), Finance, Insurance and Real Estate (H), and Services (I).

Variables include stock returns together with the Fama-French three-factor model: Market Return, SMB (Small Minus Big), and HML (High Minus Low). SMB represents the size effect between small and large firms, while HML captures the value effect based on book-to-market ratios.

3 Preliminary Analysis (EDA)

Table 2: Group statistics by the ‘industries’

Industry	Average Monthly Return	Volatility	Number of Stocks
Manufacturing	1.43%	12.87%	9
Wholesale Trade	1.23%	11.36%	6
Services	1.17%	16.48%	9
Mining	1.14%	15.95%	4
Transportation and Public Utilities	1.03%	12.25%	12
Retail Trade	0.96%	13.57%	10
Construction	0.80%	12.58%	1
Finance, Insurance and Real Estate	0.77%	7.29%	32

Table 2 summarises the performance and risk of eight industry sectors based on their monthly returns and volatility:

- **Top Performers:** Manufacturing leads with the highest average monthly return at 1.43%, followed closely by Wholesale Trade at 1.23%.
- **Risk and Volatility:** Services is the most volatile sector (16.48%), while Finance, Insurance and Real Estate is the least (7.29%).
- **Sector Representation:** The Finance sector has the largest number of stocks in this dataset (32), whereas only a single stock in Construction .

Industry Return Distributions: 2005 – 2019

Boxplots show median, quartiles, and volatility changes across industries




Figure 1: Distribution of monthly returns across industries from 2005 to 2019

As shown from the Figure 1, the industry return distributions from 2005 to 2019 reveal several key trends in market stability and sector performance:

- **Volatility Cluster from 2008 to 2009:** During this time, all industries saw a big rise in volatility, this in the taller boxes (wider interquartile ranges) and the extreme outliers, especially in Wholesale Trade, Services, and Finance.
- **Performance by sector:** Manufacturing and wholesale trade were still the best performers, with high average returns. However, manufacturing had more positive outliers over the decade.
- **Stability vs. Risk:** Finance, Insurance, and Real Estate had the most stable distributions (lowest volatility), while Mining and Services had much bigger swings.




Figure 2: Industry Correlations Heatmap

Figure 2 shows preliminary correlation analysis indicates significance systemic integration between industries. For instance, the correlation coefficient between Transportation and Public Utilities (E) and Finance (H) is 0.80, highlighting a shared sensitivity to market. Given this high dimensionality and strong inter-variable relationship, PCA is an appropriate methodology for isolating the primary component that parallel with market movements.

4 PCA Preparation

PCA variables are restricted to the 83 monthly stock return columns to prevent date or benchmark factors from biasing the extraction of common variation. Since these stocks belong to diverse sectors with various volatility, standardization is essential. Without scaling, stocks with high return variance wouldn't be proportionally dominate PCs, regardless of their true market impact. Centering and scaling ensure that all companies contribute equally. Finally, PCA is performed on this standardized matrix to transform correlated returns into linearly uncorrelated PCs.

5 PCA and Latent Factors Analysis

5.1 Q1 Does the largest principal component coincide with market returns?

To determine the number of components representing meaningful signals, we generate a Scree Plot. In high-dimensional financial data, we typically observe a sharp elbow after the first component, which indicates that one dominant latent factor is responsible for the majority of the co-movement across all companies.




Figure 3: Scree Plot of Stock Returns

	PC1	PC2	PC3	PC4	PC5
Standard deviation	4.851678	2.018468	1.666329	1.555238	1.475485
Proportion of Variance	0.283600	0.049090	0.033450	0.029140	0.026230
Cumulative Proportion	0.283600	0.332690	0.366140	0.395280	0.421510

Scree Plot of Stock Returns

The Variance Explained by the first principal component (PC1) is a critical measure of systematic risk. If PC1 accounts for a large percentage (e.g., >30%) of the total variation in an 83-stock universe, it suggests a powerful common driver.

To test if this mathematical construct is in conformity with the actual economy, we compare the PC1 scores against the observable Market Return provided in the dataset.

Correlation between PC1 and Market Return: 0.9657

A correlation magnitude near or above 0.90 indicates compelling evidence that the first principal component (PC1) is a robust proxy for the broad market. While PCA is a purely mathematical algorithm with no prior knowledge of financial indices, the fact that its primary axis of variation aligns so closely with the broad market return confirms that the market is the single most significant latent force driving individual stock returns. Consequently, PC1 can be interpreted as the Systematic Risk Factor, representing macroeconomic events that affect the entire 83-stock universe together.

5.2 Q2 Do Market Return, SMB, and HML explain common variation in stock returns?

Table 3: Correlation between Principal Components and Market Factors

component	MarketReturn	SMB	HML
PC1	0.9657	0.4817	0.2852
PC2	0.0225	0.2410	0.4157
PC3	-0.0746	-0.0082	0.0495

To examine whether market factors explain common variation in stock returns, correlations were calculated between the principal components and Market Return, SMB, and HML. As shown in Table 3, PC1 has a very strong positive correlation with Market Return (0.9657), indicating that it captures broad market movement and represents the dominant source of common variation in stock returns. PC1 also shows moderate and weak relationships with SMB (0.4817) and HML (0.2852). In contrast, PC2 and PC3 have weak relationships with Market Return, although PC2 has a moderate relationship with HML (0.4157), suggesting some value-related effects.

Table 4: Correlation between latent factor scores and market factors

Latent factor	MarketReturn	SMB	HML
Factor1	0.9865	0.4136	0.1821
Factor2	0.0754	0.2399	0.4139
Factor3	-0.0027	0.1476	0.5170

Furthermore, Table 4 provides a consistency check using the latent factor model. Factor 1 also has a very strong relationship with Market Return (0.9865), confirming that the dominant hidden structure in stock returns is overall market movement. Factor 2 and Factor 3 show moderate relationships with HML, suggesting that SMB and HML explain only smaller secondary patterns.

5.3 Q3 Which industries move in the same direction as the common factors?

Table 5: Industry-level PC1 loadings

Industry	Mean PC1 Loading	SD of Loading	Min Loading	Max Loading	Number of Stocks
Finance, Insurance and Real Estate	0.118	0.050	0.005	0.199	32
Construction	0.117	NA	0.117	0.117	1
Wholesale Trade	0.107	0.021	0.086	0.144	6
Transportation and Public Utilities	0.107	0.033	0.049	0.158	12
Retail Trade	0.098	0.024	0.054	0.131	10
Manufacturing	0.082	0.042	0.031	0.144	9
Mining	0.079	0.036	0.040	0.110	4
Services	0.065	0.022	0.037	0.101	9




Figure 4: Distribution of PC1 loadings by industry

As shown in Table 5 and Figure 4, all industries have positive PC1 loadings, indicating that they move in the same direction as the dominant common factor. This suggests that the primary latent factor reflects broad market movement rather than industry-specific behaviour.

However, the level of exposure differs across industries. Finance, Insurance and Real Estate has the highest average loading (0.118), followed by Construction (0.117) and Wholesale Trade (0.107), indicating stronger sensitivity to the common factor. In contrast, Services (0.065) and Mining (0.079) show weaker exposure.

The results also show heterogeneity within industries. Finance, Insurance and Real Estate has the largest spread of loadings, while Manufacturing and Mining also display noticeable variation. In comparison, Wholesale Trade and Retail Trade have more consistent responses. Construction contains only one stock, so heterogeneity cannot be assessed meaningfully.

Overall, industry exposure to the common factor differs substantially across sectors.

5.4 Q4 Which industry has the highest systematic risk?

Table 6: Correlation between Principal Components and Market Factors

Industry	No of stocks	Mean PC1 correlation	Mean PC1 variance share
Finance, Insurance and Real Estate	32	0.571	0.382
Construction	1	0.568	0.322
Transportation and Public Utilities	12	0.520	0.293
Wholesale Trade	6	0.521	0.280
Retail Trade	10	0.473	0.236
Manufacturing	9	0.396	0.194
Mining	4	0.384	0.171
Services	9	0.317	0.111

To identify the industry with the highest systematic risk, we use PC1 as the main market-wide factor. This is reasonable because the earlier result showed that PC1 moves very closely with the actual Market Return.

Based on the results on Table 6, Finance, Insurance and Real Estate appears to have the highest systematic risk. It has the highest average PC1 variance share, at 0.382, and the highest average PC1 correlation, at 0.571, across 32 stocks. This suggests that returns in this industry are more strongly driven by the shared market movement than the other industries.

While Construction also has a relatively high PC1 variance share, at 0.322, but this result comes from only one stock. Because of that, it is risky to say as a whole industry conclusion. Finance, Insurance and Real Estate is the stronger answer because it has both high exposure to PC1 and a much larger number of stocks. This also makes sense in practice, because financial firms are usually sensitive to wider market conditions, investor confidence, credit and economic conditions.

5.5 Q5 Which industry is counter-cyclical for hedging purposes?

Table 7: Industry exposure to PC1 for hedging assessment

Industry	No of stocks	Mean PC1 correlation	Mean PC1 variance share	Signal
Services	9	0.317	0.111	Weak positive exposure
Mining	4	0.384	0.171	Weak positive exposure

Table 7: Industry exposure to PC1 for hedging assessment

Industry	No of stocks	Mean PC1 correlation	Mean PC1 variance share	Signal
Manufacturing	9	0.396	0.194	Weak positive exposure
Retail Trade	10	0.473	0.236	Positive exposure
Transportation and Public Utilities	12	0.520	0.293	Positive exposure
Wholesale Trade	6	0.521	0.280	Positive exposure
Construction	1	0.568	0.322	Positive exposure
Finance, Insurance and Real Estate	32	0.571	0.382	Positive exposure

No industry in this sample as shown Table 7 is truly counter-cyclical because all industries have positive average PC1 correlations. This means they still tend to move in the same direction as the market wide factor. However, Services is the most defensive option because it has the lowest average PC1 correlation (0.317) and the lowest average PC1 variance share (0.111). This suggests that Services is the least exposed to broad market movements, although it should not be treated as a true hedge that moves opposite to the market.

6 Conclusion

Based on the high-dimensional analysis of NYSE stock returns from 2005 to 2019, the study concludes that **a single dominant latent factor drives common market variation**.

Key Findings and Direct Answers

Q1: Market Coincidence: The first principal component (PC1) is a robust proxy for the market, showing a 0.9657 correlation with broad market returns.

Q2: Fama-French Model: While PC1 captures overall movement, SMB and HML factors explain only a smaller portion of variation, providing only partial support for the three-factor model.

Q3: Industry Direction: All analyzed industries move in the same direction as the common factor, confirmed by universally positive PC1 loadings.

Q4: Systematic Risk: Finance, Insurance and Real Estate carries the highest systematic risk, with the highest PC1 correlation (0.571) and variance share (0.382).

Q5: Hedging: No industry is truly counter-cyclical; however, Services is the most defensive option due to its minimal market exposure.

Investment Implications

Investors should recognize that diversification across these industries does not eliminate systematic risk, as all sectors remain positively correlated with the market. For hedging, the Services sector provides the best relative protection. Conversely, the Finance sector requires careful monitoring during macroeconomic shifts due to its high sensitivity to latent market forces.

Limitations and Assumptions

This analysis assumes that linear relationships and covariance structures are sufficient to capture common variation in stock returns. PCA and factor models also assume that historical return patterns remain informative over time. However, the analysis is limited to 83 NYSE stocks and three observable market factors, meaning additional latent influences may remain unexplained. Furthermore, industry representation is uneven, with some sectors containing relatively few stocks.

7 Reference

Center for Research in Security Prices. (2026). CRSP US Stock Database [Data set]. University of Chicago Booth School of Business.

Firke S (2024). *janitor: Simple Tools for Examining and Cleaning Dirty Data*. doi:10.32614/CRAN.package.janitor <https://doi.org/10.32614/CRAN.package.janitor>, R package version 2.2.1, <https://CRAN.R-project.org/package=janitor>.

Wickham H, Averick M, Bryan J, Chang W, McGowan LD, François R, Golemund G, Hayes A, Henry L, Hester J, Kuhn M, Pedersen TL, Miller E, Bache SM, Müller K, Ooms J, Robinson D, Seidel DP, Spinu V, Takahashi K, Vaughan D, Wilke C, Woo K, Yutani H (2019). “Welcome to the tidyverse.” *Journal of Open Source Software*, 4(43), 1686. doi:10.21105/joss.01686 <https://doi.org/10.21105/joss.01686>.

Wickham H (2025). *stringr: Simple, Consistent Wrappers for Common String Operations*. doi:10.32614/CRAN.package.stringr <https://doi.org/10.32614/CRAN.package.stringr>, R package version 1.6.0, <https://CRAN.R-project.org/package=stringr>.

Wickham H, Pedersen T, Seidel D (2025). *scales: Scale Functions for Visualization*. doi:10.32614/CRAN.package.scales <https://doi.org/10.32614/CRAN.package.scales>, R package version 1.4.0, <https://CRAN.R-project.org/package=scales>.

William Revelle (2026). *psych: Procedures for Psychological, Psychometric, and Personality Research*. Northwestern University, Evanston, Illinois. R package version 2.6.3, <https://CRAN.R-project.org/package=psych>.

Xie Y (2025). *knitr: A General-Purpose Package for Dynamic Report Generation in R*. R package version 1.51, <https://yihui.org/knitr/>.

7.1 AI Usage Statement

Generative AI was used to assist with selected parts of this assignment, mainly to improve coding efficiency, refine wording and grammar, and support the interpretation of several analytical concepts. AI suggestions were used to help produce cleaner and more efficient code, improve the clarity of explanations, and assist with interpreting PCA and factor models in relation to market movement. In addition, AI was used to support the analysis in understanding the relationship between principal components, Fama-French factors, industry movement, and the discussion of counter-cyclical industries. The main ideas, analytical decisions, coding implementation, interpretation, and final version of the assignment were still completed and reviewed by us.

- Link 1
- Link 2
- Link 3

8 Appendix

8.1 Data and Preparation Code

```
# 1. Load Libraries
library(tidyverse)
library(janitor)
library(psych)
library(knitr)
library(stringr)
library(scales)

# 2. Load Data
stocks <- read_csv("SampleH.csv")
market_factors <- read_csv("MarketFactors.csv")

# 3. Define the mapping of letters to Industry Names
industry_lookup <- c(
  "B" = "Mining",
  "C" = "Construction",
  "D" = "Manufacturing",
```

```

"E" = "Transportation and Public Utilities",
"F" = "Wholesale Trade",
"G" = "Retail Trade",
"H" = "Finance, Insurance and Real Estate",
"I" = "Services"
)

# 4. Transform and Merge Data
# We pivot the stocks to a 'long' format so each row is a unique Stock-Date observation
data_complete <- stocks |>
  pivot_longer(
    # Keep Date as is
    cols = -Date,
    # New column for stock codes
    names_to = "stock_id",
    # new column for the return values
    values_to = "return"
  ) |>
  mutate(industry_letter = substr(stock_id, 1, 1)) |>
  mutate(industry_name = industry_lookup[industry_letter]) |>
  inner_join(market_factors, by = "Date") |>
  mutate(year = substr(Date, 2, 5)) |>
  # Reorder: put `Return` column after `industry_name`
  relocate(return, .after = industry_name) |>
  relocate(year, .before = Date) |>
  # clean names
  clean_names()

kable(
  head(
    data_complete |>
      select(
        stock_id,
        industry_name,
        return,
        market_return,
        smb,
        hml
      )
  ),
  col.names = c(
    "Stock ID",

```

```

    "Industry",
    "Return",
    "Market Return",
    "SMB",
    "HML"
  ),
  digits = 3
)

```

8.2 Preliminary Analysis Code

Table group statistics by the 'industries'

```

# 1. Group statistics by the 'industries' vector created earlier
eda_summary <- data_complete |>
  group_by(industry_name) |>
  summarise(
    avg_return_monthly = label_percent(accuracy = 0.01)(mean(return, na.rm = TRUE)),
    volatility = label_percent(accuracy = 0.01)(sd(return, na.rm = TRUE)),
    num_stocks = n_distinct(stock_id)
  ) |>
  arrange(desc(avg_return_monthly))

eda_summary_clean <- eda_summary |>
  rename(
    Industry = industry_name,
    `Average Monthly Return` = avg_return_monthly,
    Volatility = volatility,
    `Number of Stocks` = num_stocks
  )

```

Time-series boxplots

```

# Create combined faceted boxplot
data_complete |>
  mutate(
    industry_name = factor(
      industry_name,
      levels = industry_order
    )
  ) |>

```

```

ggplot(
  aes(
    x = factor(year),
    y = return,
    fill = industry_name
  )
) +
geom_boxplot(
  outlier.size = 0.5,
  alpha = 0.7,
  width = 0.5,
  median.linewidth = 0.8
) +
facet_wrap(
  ~ industry_name,
  scales = "free_y",
  ncol = 4,
  labeller = label_wrap_gen(width = 20)
) +
geom_hline(
  yintercept = 0,
  linetype = "dashed",
  color = "gray50",
  alpha = 0.5
) +
scale_x_discrete(
  breaks = seq(2005, 2019, by = 2)
) +
labs(
  title = "Industry Return Distributions: 2005 - 2019",
  subtitle = "Boxplots show median, quartiles, and volatility changes across industries",
  x = "Year",
  y = "Monthly Return (%)"
) +
theme_minimal(base_size = 11) +
theme(
  plot.title = element_text(face = "bold"),
  strip.text = element_text(face = "bold", size = 9),
  axis.text.x = element_text(
    angle = 45,
    hjust = 1,
    size = 7
  )
)

```



```

),
axis.text.y = element_text(size = 8),
panel.spacing = unit(1, "lines"),
legend.position = "none"
)

```

Heatmap plot

```

# 3. Calculate Average Monthly Returns per Industry (time series) for correlation analysis
industry_ts <- data_complete |>
  # Group by Industry and Date
  group_by(industry_name, date) |>
  # Calculate the average return for the each industry classification for that month
  summarise(avg_return = mean(return, na.rm = TRUE), .groups = "drop") |>
  # Pivot so each industry gets its own column
  pivot_wider(names_from = industry_name, values_from = avg_return)

# 4. Create Correlation Matrix between Industry Averages
ind_cor_long <- industry_ts |>
  select(-date) |>
  cor(use = "complete.obs") |>
  as.table() |>
  as.data.frame()

# 5. Use heatmap to visualise
ggplot(ind_cor_long, aes(Var1, Var2, fill = Freq)) +
  geom_tile(color = "white") +
  geom_text(aes(label = round(Freq, 2)),
            color = ifelse(abs(ind_cor_long$Freq) > 0.7, "white", "black"), size = 2.5) +
  scale_fill_gradient2(low = "#3B4CC0", high = "#B40427", mid = "white",
                      midpoint = 0, limit = c(-1,1), name="Correlation") +
  coord_fixed() +
  theme_minimal() +
  labs(title = "Industry Correlations Heatmap", x = NULL, y = NULL) +
  scale_x_discrete(labels = function(x) str_trunc(x, 15)) +
  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust = 1),
        panel.grid.major = element_blank(), # Remove background grid
        legend.position = "right"
  )

```

8.3 PCA Preparation Code

```
# 1. Select only stock return columns (numeric variables)
# We exclude 'date' and the Fama-French factors
pca_input <- stocks |> select(-Date)

# 2. Standardise and Run PCA
# 'scale. = TRUE' ensures data is standardised (mean=0, sd=1)
pca_res <- prcomp(pca_input, scale. = TRUE)
```

8.4 Q1 Code

Scree Plot

```
# 1. Visualize Eigenvalues
screeplot(pca_res, type = "lines", main = "Scree Plot of Stock Returns")
abline(h = 1, col="red", lty=2) # Kaiser Criterion threshold

# 2. Calculate Variance Explained
kable(summary(pca_res)$importance[, 1:5])
```

Comparison PC1 with Market Return

```
# 1. Extract PC1 scores
pc1_scores <- pca_res$x[, 1]

# 2. Correlate with Market Return
market_test_corr <- cor(pc1_scores, market_factors$MarketReturn)

cat("Correlation between PC1 and Market Return:", round(market_test_corr, 4))
```

8.5 Q2 Code

Table: Correlation between Principal Components and Market Factors

```
# Extract PCA scores
pc_scores <- as.data.frame(pca_res$x) |>
  mutate(Date = stocks$Date)
```

```

# Join PCA scores with market factors by Date
q2_data <- pc_scores |>
  inner_join(market_factors, by = "Date")

# Correlation table between first 3 PCs and market factors
q2_cor_table <- q2_data |>
  select(PC1, PC2, PC3, MarketReturn, SMB, HML) |>
  cor(use = "complete.obs") |>
  as.data.frame() |>
  rownames_to_column("component") |>
  filter(component %in% c("PC1", "PC2", "PC3")) |>
  select(component, MarketReturn, SMB, HML) |>
  mutate(across(where(is.numeric), ~ round(.x, 4)))

kable(q2_cor_table)

```

Table: Factor Correlation Analysis

```

factor_res <- factanal(
  x = pca_input,
  factors = 3,
  scores = "regression",
  rotation = "none",
  lower = 0.05
)

factor_scores <- as_tibble(factor_res$scores) |>
  set_names(paste0("Factor", 1:3)) |>
  mutate(Date = stocks$Date)

# Factor signs are arbitrary, just like PCA signs.
# This aligns Factor 1 so positive values mean movement with Market Return.
factor1_sign <- ifelse(
  cor(factor_scores$Factor1, market_factors$MarketReturn, use = "complete.obs") < 0,
  -1,
  1
)

factor_scores <- factor_scores |>
  mutate(Factor1 = factor1_sign * Factor1)

factor_market_cor <- factor_scores |>

```

```

inner_join(market_factors, by = "Date") |>
select(Factor1, Factor2, Factor3, MarketReturn, SMB, HML) |>
cor(use = "complete.obs") |>
as.data.frame() |>
rownames_to_column("factor") |>
filter(factor %in% c("Factor1", "Factor2", "Factor3")) |>
select(factor, MarketReturn, SMB, HML) |>
mutate(across(where(is.numeric), ~ round(.x, 4))) |>
rename(
  `Latent factor` = factor
)

kable(factor_market_cor)

```

8.6 Q3 Code

Table: Industry-level PC1 loadings

```

pc1_loadings |>
ggplot(aes(x = reorder(industry, pc1_loading, median), y = pc1_loading)) +
geom_boxplot(fill = "lightblue", color = "black", width = 0.4, size = 0.7) +
coord_flip() +
labs(
  title = "PC1 Loadings by Industry",
  x = "Industry",
  y = "PC1 Loading"
) +
theme_minimal(base_size = 12) +
theme(
  plot.title = element_text(face = "bold"),
  axis.text.y = element_text(size = 10),
  axis.text.x = element_text(size = 10)
)

```

Boxplot: Distribution of PC1 loadings by industry

```

pc1_loadings |>
ggplot(aes(x = reorder(industry, pc1_loading, median), y = pc1_loading)) +
geom_boxplot(fill = "lightblue", color = "black", width = 0.4, size = 0.7) +
coord_flip() +
labs(

```

```

    title = "PC1 Loadings by Industry",
    x = "Industry",
    y = "PC1 Loading"
  ) +
  theme_minimal(base_size = 12) +
  theme(
    plot.title = element_text(face = "bold"),
    axis.text.y = element_text(size = 10),
    axis.text.x = element_text(size = 10)
  )

```

8.7 Q4 Code

Industry risk based on PC1

```

pc1_sign <- ifelse(
  cor(pca_res$x[, 1], market_factors$MarketReturn, use = "complete.obs") < 0,
  -1,
  1
)

# Stock exposure to PC1
stock_pc1 <- tibble(
  stock_id = colnames(pca_input),
  industry_letter = substr(stock_id, 1, 1),
  industry_name = as.character(industry_lookup[industry_letter]),
  pc1_loading = pc1_sign * pca_res$rotation[, 1],
  pc1_corr = pc1_loading * pca_res$sdev[1],
  pc1_variance_share = pc1_corr^2
)

q4_industry_risk <- stock_pc1 |>
  group_by(industry_name) |>
  summarise(
    n_stocks = n(),
    mean_pc1_corr = mean(pc1_corr, na.rm = TRUE),
    mean_pc1_variance_share = mean(pc1_variance_share, na.rm = TRUE),
    .groups = "drop"
  ) |>
  arrange(desc(mean_pc1_variance_share)) |>
  mutate(across(-c(industry_name, n_stocks), ~ round(.x, 3))) |>

```

```

  rename(
    Industry = industry_name,
    `No of stocks` = n_stocks,
    `Mean PC1 correlation` = mean_pc1_corr,
    `Mean PC1 variance share` = mean_pc1_variance_share
  )

kable(q4_industry_risk)

```

8.8 Q5 Code

Industry exposure to PC1

```

q5_hedging <- stock_pc1|>
  group_by(industry_name)|>
  summarise(
    n_stocks = n(),
    mean_pc1_corr = mean(pc1_corr, na.rm = TRUE),
    mean_pc1_variance_share = mean(pc1_variance_share, na.rm = TRUE),
    .groups = "drop"
  )|>
  arrange(mean_pc1_corr)|>
  mutate(
    signal = case_when(
      mean_pc1_corr < 0 ~ "Counter-cyclical",
      mean_pc1_corr < 0.40 ~ "Weak positive exposure",
      TRUE ~ "Positive exposure"
    )
  )|>
  mutate(across(-c(industry_name, n_stocks, signal), ~ round(.x, 3)))|>
  rename(
    Industry = industry_name,
    `No of stocks` = n_stocks,
    `Mean PC1 correlation` = mean_pc1_corr,
    `Mean PC1 variance share` = mean_pc1_variance_share,
    `Signal` = signal
  )

kable(q5_hedging)

```