

High-Dimensional Data Analysis Project

Group 3H

2026-04-01

Table of contents

Introduction (120 words)	2
Data and Preparation (150 words)	2
Preliminary Analysis (EDA) (180 words)	3
PCA Preparation (100 words)	5
PCA and Latent Factors Analysis (800 words)	5
Q1 Does the largest principal component coincide with market returns?	5
Q2 Do Market Return, SMB, and HML explain common variation in stock returns?	7
Q3 Which industries move in the same direction as the common factors?	7
Q4 Which industry has the highest systematic risk?	9
Q5 Which industry is counter-cyclical for hedging purposes?	10
Conclusion (150 words)	10
Reference	10
Appendix	10
Data and Preparation Code	10
Preliminary Analysis Code	10
PCA Preparation Code	10
Q1 Code	10
Q2 Code	10
Q3 Code	11
Q4 Code	12
Q5 Code	12

Introduction (120 words)

The analysis of stock returns is a key task in finance, aimed at understanding the relationship between risk and return. This report analyses the monthly returns of 100 New York Stock Exchange listed companies from 2005 to 2019. Using high-dimensional data techniques, the study separates stock returns into systematic risk driven by overall market movements and idiosyncratic risk specific to individual firms or industries.

The objective is to examine whether latent factors align with observable market benchmarks such as the S&P 500. In particular, the analysis tests whether the first principal component represents the market, whether the Fama-French three-factor model explains common variation, which industries move with these factors, which industries carry higher systematic risk, and which industries may provide hedging benefits.

Data and Preparation (150 words)

- Important variables description
- Data cleaning and formatting
- Industry classification from stock identifiers
- Data validation (dimensions and missing values)

Data preparation is the foundation of high-dimensional analysis. We utilize `SampleH.csv` (stock returns) and `MarketFactors.csv` (benchmarks).

```
# A tibble: 6 x 9
  year date stock_id industry_letter industry_name return market_return smb
  <chr> <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl>
1 2005 Y200~ F85464 F Wholesale Tr~ -0.0280 -2.76 -
1.66
2 2005 Y200~ H89060 H Finance, Ins~ -0.0694 -2.76 -
1.66
3 2005 Y200~ H87476 H Finance, Ins~ -0.0402 -2.76 -
1.66
4 2005 Y200~ E79879 E Transportati~ -0.322 -2.76 -
1.66
5 2005 Y200~ H10777 H Finance, Ins~ -0.0422 -2.76 -
1.66
6 2005 Y200~ B89901 B Mining 0.154 -2.76 -
1.66
# i 1 more variable: hml <dbl>
```

```
[1] 14940      9
```

```
[1] 0
```

The final dataset contains 14,940 observations across 9 variables. This reflects the longitudinal nature of the data, covering multiple stocks over 180 months (January 2005 - December 2019). The data likely contained 83 stocks ($14,940 / 180 \text{ months} = 83$).

The industry classification maps stocks to sectors such as Mining (B), Construction (C), Manufacturing (D), Transportation and Public Utilities (E), Wholesale Trade (F), Retail Trade (G), Finance, Insurance and Real Estate (H) and Services (I).

Variables include stock returns and three factors: Market Return, SMB (Small Minus Big, the size effect factor, capturing compensation for investing in small capitalization companies relative to large capitalization companies), and HML (High Minus Low, the value factor, capturing the compensation for value as measured by the book to market ratio).


Preliminary Analysis (EDA) (180 words)

- Summary statistic by Industry (mean, sd, no. of stocks)
- Time series analysis of industry AVG returns
- Correlation structure between industries


```
# A tibble: 8 x 4
  industry_name      avg_return volatility num_stocks
  <chr>            <dbl>      <dbl>      <int>
1 Construction      0.00796    0.126         1
2 Finance, Insurance and Real Estate 0.00766    0.0729        32
3 Manufacturing      0.0143     0.129         9
4 Mining             0.0114     0.159         4
5 Retail Trade       0.00960    0.136        10
6 Services           0.0117     0.165         9
7 Transportation and Public Utilities 0.0103     0.122        12
8 Wholesale Trade    0.0123     0.114         6
```

Industry Return Distributions: 2005 – 2019

Boxplots show median, quartiles, and volatility shifts per year



Industry Correlations



Preliminary correlation analysis indicates significant systemic integration between industries. For instance, the correlation coefficient between sectors Transportation and Public Utilities (E) and Finance, Insurance and Real Estate (H) is 0.80, highlighting a shared sensitivity to market conditions. Given this high dimensionality and strong inter-variable relationship, PCA is an

appropriate methodology for isolating the primary component that parallel with broad market movements.

PCA Preparation (100 words)

- Selection of stocks return variables
- Data standardisation
- PCA implementation

The selection of variables for PCA is limited to the 83 monthly stock return columns, as including the date or benchmark factors would bias the mathematical extraction of common variation. Since these stocks belong to diverse sectors with varying levels of price volatility, standardisation is a must. Without scaling, stocks with the highest return variance would disproportionately influence the principal components, regardless of their actual contribution to market-wide trends. By centering and scaling the data, we ensure each company contributes equally to the analysis. The PCA is then implemented on this standardised matrix to transform the correlated returns into a set of linearly uncorrelated principal components.


PCA and Latent Factors Analysis (800 words)

Q1 Does the largest principal component coincide with market returns?

- Scree plot and component selection
- Variance explained by principal components
- Comparison of PC1 with Market Return
- Interpretation of PC1 as a market factor

To determine the number of components representing meaningful signals, we generate a Scree Plot. In high-dimensional financial data, we typically observe a sharp elbow after the first component, which indicates that one dominant latent factor is responsible for the majority of the co-movement across all companies.

Scree Plot of Stock Returns



	PC1	PC2	PC3	PC4	PC5
Standard deviation	4.851678	2.018468	1.666329	1.555238	1.475485
Proportion of Variance	0.283600	0.049090	0.033450	0.029140	0.026230
Cumulative Proportion	0.283600	0.332690	0.366140	0.395280	0.421510

The Variance Explained by the first principal component (PC1) is a critical measure of systematic risk. If PC1 accounts for a large percentage (e.g., >30%) of the total variation in an 83-stock universe, it suggests a powerful common driver.

To test if this mathematical construct is in conformity with the actual economy, we compare the PC1 scores against the observable Market Return provided in the dataset.

Correlation between PC1 and Market Return: 0.9657

A correlation magnitude near or above 0.90 indicates compelling evidence that the first principal component (PC1) is a robust proxy for the broad market. While PCA is a purely mathematical algorithm with no prior knowledge of financial indices, the fact that its primary axis of variation aligns so closely with the broad market return confirms that the market is the single most significant latent force driving individual stock returns. Consequently, PC1 can be interpreted as the Systematic Risk Factor, representing macroeconomic events that affect the entire 83-stock universe together.

Q2 Do Market Return, SMB, and HML explain common variation in stock returns?

Table 1: Correlation between Principal Components and Market Factors

component	MarketReturn	SMB	HML
PC1	0.9657	0.4817	0.2852
PC2	0.0225	0.2410	0.4157
PC3	-0.0746	-0.0082	0.0495

To examine whether market factors explain common variation in stock returns, we combined the principal component scores with the market factors using the date variable. We then calculated correlations between the principal components and Market Return, SMB, and HML.

As shown in Table 1, PC1 has a very strong positive relationship with Market Return (0.9657), which means it captures overall market movement. PC1 also has a moderate relationship with SMB (0.4817) and a weaker relationship with HML (0.2852). PC2 has very little relationship with Market Return (0.0225), but shows a moderate relationship with HML (0.4157), suggesting it may capture value-related effects. PC3 shows very weak relationships with all three factors.

These results show that Market Return explains most of the variation in stock returns, while SMB and HML only explain a smaller part. This means the Fama-French three-factor model is only partially supported and does not fully explain the patterns found in the data.

Q3 Which industries move in the same direction as the common factors?

Table 2: Industry-level PC1 loadings

Industry	Mean PC1 Loading	SD of Loading	Min Loading	Max Loading	Number of Stocks
Finance, Insurance and Real Estate	0.118	0.050	0.005	0.199	32
Construction	0.117	NA	0.117	0.117	1
Wholesale Trade	0.107	0.021	0.086	0.144	6
Transportation and Public Utilities	0.107	0.033	0.049	0.158	12
Retail Trade	0.098	0.024	0.054	0.131	10
Manufacturing	0.082	0.042	0.031	0.144	9
Mining	0.079	0.036	0.040	0.110	4
Services	0.065	0.022	0.037	0.101	9




Figure 1: Distribution of PC1 loadings by industry

As shown in Table 2 and Figure 1, all industries have positive PC1 loadings, indicating that they move in the same direction as the dominant common factor. This confirms that the common factor affects the entire market rather than a specific group of industries.

However, the strength of this relationship differs across industries. Finance, Insurance and Real Estate have the highest average loading (0.118), followed closely by Construction (0.117) and Wholesale Trade (0.107), suggesting these industries are more sensitive to the common

factor. In contrast, Services (0.065) and Mining (0.079) have lower average loadings, indicating weaker exposure.

More importantly, there is clear heterogeneity within industries. Finance, Insurance and Real Estate shows a relatively large spread ($SD = 0.050$ and a wide range from 0.005 to 0.199), indicating that stocks within this industry respond very differently to the common factor. Manufacturing and Mining also display noticeable variation, suggesting inconsistent behaviour across stocks. In contrast, Wholesale Trade and Retail Trade show smaller spreads, indicating more homogeneous responses.

The Construction industry appears as a single line in Figure 1 and has no variation in Table 2 because it contains only one stock. As a result, it is not meaningful to assess heterogeneity within this industry.

Overall, while all industries move in the same direction as the common factor, the key result is that the degree of heterogeneity varies substantially across industries, suggesting that industry classification alone does not fully explain stock behaviour.

Q4 Which industry has the highest systematic risk?

- Measurement of systematic risk using factor loadings
- Comparison of absolute loadings across industries
- Identification of high-risk industries

Table 3: Correlation between Principal Components and Market Factors

Industry	No of stocks	Mean PC1 correlation	Mean PC1 variance share
Finance, Insurance and Real Estate	32	0.571	0.382
Construction	1	0.568	0.322
Transportation and Public Utilities	12	0.520	0.293
Wholesale Trade	6	0.521	0.280
Retail Trade	10	0.473	0.236
Manufacturing	9	0.396	0.194
Mining	4	0.384	0.171
Services	9	0.317	0.111

To identify the industry with the highest systematic risk, we use PC1 as the main market-wide factor. This is reasonable because the earlier result showed that PC1 moves very closely with the actual Market Return.

Based on the results on Table 3, Finance, Insurance and Real Estate appears to have the highest systematic risk. It has the highest average PC1 variance share, at 0.382, and the highest average PC1 correlation, at 0.571, across 32 stocks. This suggests that returns in this industry are more strongly driven by the shared market movement than the other industries.

While Construction also has a relatively high PC1 variance share, at 0.322, but this result comes from only one stock. Because of that, it is risky to say as a whole industry conclusion. Finance, Insurance and Real Estate is the stronger answer because it has both high exposure to PC1 and a much larger number of stocks. This also makes sense in practice, because financial firms are usually sensitive to wider market conditions, investor confidence, credit and economic conditions.

Q5 Which industry is counter-cyclical for hedging purposes?

- Identification of negative or weak loadings
- Industry behaviour relative to market factor
- Hedging implications

Conclusion (150 words)

- Summary of key findings
- Direct answers to Q1 - Q5
- Implications for investment decisions

Reference

Appendix

Data and Preparation Code

Preliminary Analysis Code

PCA Preparation Code

Q1 Code

Q2 Code

Table: Correlation between Principal Components and Market Factors

```

# Extract PCA scores
pc_scores <- as.data.frame(pca_res$x) |>
  mutate(Date = stocks$Date)

# Join PCA scores with market factors by Date
q2_data <- pc_scores |>
  inner_join(market_factors, by = "Date")

# Correlation table between first 3 PCs and market factors
q2_cor_table <- q2_data |>
  select(PC1, PC2, PC3, MarketReturn, SMB, HML) |>
  cor(use = "complete.obs") |>
  as.data.frame() |>
  rownames_to_column("component") |>
  filter(component %in% c("PC1", "PC2", "PC3")) |>
  select(component, MarketReturn, SMB, HML) |>
  mutate(across(where(is.numeric), ~ round(.x, 4)))

kable(q2_cor_table)

```

Q3 Code

Table: Industry-level PC1 loadings

```

pc1_loadings |>
  ggplot(aes(x = reorder(industry, pc1_loading, median), y = pc1_loading)) +
  geom_boxplot(fill = "lightblue", color = "black", width = 0.4, size = 0.7) +
  coord_flip() +
  labs(
    title = "PC1 Loadings by Industry",
    x = "Industry",
    y = "PC1 Loading"
  ) +
  theme_minimal(base_size = 12) +
  theme(
    plot.title = element_text(face = "bold"),
    axis.text.y = element_text(size = 10),
    axis.text.x = element_text(size = 10)
  )

```

Boxplot: Distribution of PC1 loadings by industry

```
pc1_loadings |>
  ggplot(aes(x = reorder(industry, pc1_loading, median), y = pc1_loading)) +
  geom_boxplot(fill = "lightblue", color = "black", width = 0.4, size = 0.7) +
  coord_flip() +
  labs(
    title = "PC1 Loadings by Industry",
    x = "Industry",
    y = "PC1 Loading"
  ) +
  theme_minimal(base_size = 12) +
  theme(
    plot.title = element_text(face = "bold"),
    axis.text.y = element_text(size = 10),
    axis.text.x = element_text(size = 10)
  )
```

Q4 Code

Q5 Code